

Lab8

和死锁斗争到死的lab，整个lab在做的就是把一个大锁拆成几个小锁

Memory allocator

这个lab任务是将整个共享内存拆成每个cpu一人一个共享内存，然后每个cpu一人一把锁，以此实现优化

首先把内存拆成N个内存：

```
// 为每个CPU设置freelist
struct
{
    struct spinlock lock;
    struct run *freelist;
} kmem[NCPU];
```

然后修改kinit，能够正确初始化所有的锁。空闲空间暂时放在一个cpu上，之后再处理

根据提示，把所有锁都命名为kmem也没事

```
void kinit()
{
    //初始化锁
    for (int i = 0; i < NCPU; i++)
    {
        initlock(&kmem[i].lock, "kmem");
    }

    freerange(end, (void *)PHYSTOP);
}
```

修改kfree

注意，获取cpuid需要关闭中断

```
void kfree(void *pa)
{
    struct run *r;

    if (((uint64)pa % PGSIZE) != 0 || (char *)pa < end || (uint64)pa >= PHYSTOP)
        panic("kfree");

    // Fill with junk to catch dangling refs.
    memset(pa, 1, PGSIZE);

    r = (struct run *)pa;

    //关闭中断以获取正确的CPUid
```

```

push_off();
int id = cpuid();
pop_off();

acquire(&kmem[id].lock);
r->next = kmem[id].freelist;
kmem[id].freelist = r;
release(&kmem[id].lock);
}

```

修改kalloc

如果自己空间不够了就去其他cpu上偷一块空间回来

```

void *
kalloc(void)
{
    struct run *r;
    //关闭中断以获取正确的CPUid
    push_off();
    int id = cpuid();
    pop_off();

    acquire(&kmem[id].lock);
    r = kmem[id].freelist;
    if (r)
        kmem[id].freelist = r->next;
    else{
        //去其他cpu上偷一块freelist
        for(int i=0;i<NCPU;i++){
            //跳过自己
            if(id==i)
                continue;
            //上锁访问
            acquire(&kmem[i].lock);
            r=kmem[i].freelist;
            //偷
            if(r){
                kmem[i].freelist=kmem[i].freelist->next;
            }
            release(&kmem[i].lock);
            //偷到了
            if(r){
                break;
            }
            //在释放的时候，这块地址就会被插入到当前CPU的freelist中
        }
    }
    release(&kmem[id].lock);

    if (r)
        memset((char *)r, 5, PGSIZE); // fill with junk
    return (void *)r;
}

```

```
}
```

Buffer Cache

任务是减小buffer粒度，通过哈希把cache拆成几个，一人一个锁。这样就避免多线程竞争共享cache实现性能提高

修改在bio里

根据提示，需要大锁加小锁避免死锁：

```
struct
{

    struct buf buf[NBUF];

    // 拆成几个bucket
    struct spinlock lock[NBUCKET];
    struct buf head[NBUCKET];

    // 大锁
    struct spinlock biglock;

} bcache;
```

同时需要修改buf结构体，通过时间间隔来实现LRU

```
struct buf {
    int valid;    // has data been read from disk?
    int disk;    // does disk "own" buf?
    uint dev;
    uint blockno;
    struct sleeplock lock;
    uint refcnt;
    struct buf *prev; // LRU cache list
    struct buf *next;
    uchar data[BSIZE];
    uint timestamp;
};

#define NBUCKET 13
```

初始化锁和链表

```
void binit(void)
{
    struct buf *b;

    initlock(&bcache.biglock, "biglock");
    for (int i = 0; i < NBUCKET; i++)
    {
```

```

        initlock(&bcache.lock[i], "lock");
    }

    for (int i = 0; i < NBUCKET; i++)
    {
        bcache.head[i].next = &bcache.head[i];
        bcache.head[i].prev = &bcache.head[i];
    }

    // 全放到第0个
    for (b = bcache.buf; b < bcache.buf + NBUF; b++)
    {
        b->next = bcache.head[0].next;
        b->prev = &bcache.head[0];
        initsleeplock(&b->lock, "buffer");
        bcache.head[0].next->prev = b;
        bcache.head[0].next = b;
    }
}

```

修改brelse, bpin和bunpin, 改成数组

```

void brelse(struct buf *b)
{
    if (!holdingsleep(&b->lock))
        panic("brelse");

    releasesleep(&b->lock);

    int bucketkey = hash(b->blockno);

    acquire(&bcache.lock[bucketkey]);
    b->refcnt--;
    if (b->refcnt == 0)
    {
        // no one is waiting for it.
        // b->next->prev = b->prev;
        // b->prev->next = b->next;
        // b->next = bcache.head.next;
        // b->prev = &bcache.head;
        // bcache.head.next->prev = b;
        // bcache.head.next = b;

        // 由于没有人用, 设置最后使用时间
        b->timestamp = ticks;
    }

    release(&bcache.lock[bucketkey]);
}

void bpin(struct buf *b)
{
    int bucketkey = hash(b->blockno);
}

```

```

    acquire(&bcache.lock[bucketkey]);
    b->refcnt++;
    release(&bcache.lock[bucketkey]);
}

void bunpin(struct buf *b)
{
    int bucketkey = hash(b->blockno);
    acquire(&bcache.lock[bucketkey]);
    b->refcnt--;
    release(&bcache.lock[bucketkey]);
}

```

在brelse维护timestamp，记录最后一次被使用的时间来实现LRU

修改bget，步骤如下：

- 命中cache，无事发生，返回即可
- 没命中cache，准备存入cache。需要释放小锁锁大锁（为了后面的偷内存）
- 释放小锁和上大锁的中间可能cache更新，重新检查一遍当前cpu内内存（避免了散列到旧块相同的位置）
- 还是没命中，就在当前bucket内根据timestamp找到最后一次使用间隔最久的空闲块，存储
- 没有空闲块，去其他bucket中偷一块来

```

static struct buf *
bget(uint dev, uint blockno)
{
    struct buf *b;
    int time = 0;
    struct buf *SelectFreeBlock = 0;
    int hashkey = hash(blockno);
    acquire(&bcache.lock[hashkey]);

    // Is the block already cached?
    // 如果cached了，直接返回
    for (b = bcache.head[hashkey].next; b != &bcache.head[hashkey]; b = b->next)
    {
        if (b->dev == dev && b->blockno == blockno)
        {
            b->refcnt++;
            release(&bcache.lock[hashkey]);
            acquiresleep(&b->lock);
            return b;
        }
    }
    release(&bcache.lock[hashkey]);

    // Not cached.
    // 避免死锁，先大锁稳住再上小锁
    acquire(&bcache.biglock);

```

```
// 1. 检查在重新上锁的时候当前bucket有没有更新
acquire(&bcache.lock[hashkey]);
for (b = bcache.head[hashkey].next; b != &bcache.head[hashkey]; b = b->next)
{
    if (b->dev == dev && b->blockno == blockno)
    {
        b->refcnt++;
        release(&bcache.lock[hashkey]);
        release(&bcache.biglock);
        acquiresleep(&b->lock);
        return b;
    }
}
```

// 2. 在当前位置根据时间戳找个空闲块返回回去

```
for (b = bcache.head[hashkey].next; b != &bcache.head[hashkey]; b = b->next)
{
    if (b->refcnt == 0 && (SelectFreeBlock == 0 || b->timestamp < time))
    {
        SelectFreeBlock = b;
        time = b->timestamp;
    }
}
```

// 找得到就返回

```
if (SelectFreeBlock)
{
    SelectFreeBlock->refcnt++;
    SelectFreeBlock->dev = dev;
    SelectFreeBlock->valid = 0;
    SelectFreeBlock->blockno = blockno;

    release(&bcache.lock[hashkey]);
    release(&bcache.biglock);
    acquiresleep(&SelectFreeBlock->lock);
    return SelectFreeBlock;
}
```

// 3. 从其他bucket中找一个block返回回去

```
for (int tmp = hash(hashkey + 1); tmp != hashkey; tmp = hash(tmp + 1))
{
    acquire(&bcache.lock[tmp]);
    for (b = bcache.head[tmp].next; b != &bcache.head[tmp]; b = b->next)
    {
        if (b->refcnt == 0 && (SelectFreeBlock == 0 || b->timestamp < time))
        {
            SelectFreeBlock = b;
            time = b->timestamp;
        }
    }
}

if (SelectFreeBlock)
```

```

{
    SelectFreeBlock->refcnt++;
    SelectFreeBlock->dev = dev;
    SelectFreeBlock->valid = 0;
    SelectFreeBlock->blockno = blockno;

    // 把这个内存块放到指定哈希队列表首
    SelectFreeBlock->next->prev = SelectFreeBlock->prev;
    SelectFreeBlock->prev->next = SelectFreeBlock->next;
    release(&bcache.lock[tmp]);
    SelectFreeBlock->next = bcache.head[hashkey].next;
    SelectFreeBlock->prev = &bcache.head[hashkey];
    bcache.head[hashkey].next->prev = SelectFreeBlock;
    bcache.head[hashkey].next = SelectFreeBlock;

    release(&bcache.lock[hashkey]);
    release(&bcache.biglock);
    acquiresleep(&SelectFreeBlock->lock);
    return SelectFreeBlock;
}

release(&bcache.lock[tmp]);
}

release(&bcache.lock[hashkey]);
release(&bcache.biglock);
panic("bget: no buffers");
}

```

实验结果

```

== Test running kalloc test == (77.5s)
== Test   kalloc test: test1 ==
    kalloc test: test1: OK
== Test   kalloc test: test2 ==
    kalloc test: test2: OK
== Test kalloc test: sbrkmuch == kalloc test: sbrkmuch: OK (10.4s)
== Test running bcachetest == (8.8s)
== Test   bcachetest: test0 ==
    bcachetest: test0: OK
== Test   bcachetest: test1 ==
    bcachetest: test1: OK
== Test usertests == usertests: OK (136.2s)

```

实验小结

和死锁斗争的lab。最痛苦的地方在于代码抄错变量了结果锁会panic（死锁或者typo error），查了很久还以为是锁上的有问题，最后重新读了一遍代码发现是双向链表写错了.....

最大的收获是理解怎么解决死锁了。如果每个资源都需要两个资源但可能只拿到一个资源导致死锁，那么就
让拿两个资源这个行动变成原子化操作，就不会出现自旋竞争锁了（大锁-小锁拓扑排序确保上锁-解锁是按
顺序来的）