

# Lab10

感觉是最麻烦的lab了

整个lab就是在实现mmap这个系统调用

## mmap

---

首先一通操作把mmap和munmap系统调用注册上去，和lab2类似不再赘述

然后开始实现mmap

- 抄课本，抄一个vma结构体声明，并修改proc

```
#define VMASIZE 16

struct vma {
    int used;
    uint64 addr;
    int length;
    int prot;
    int flags;
    int fd;
    int offset;
    struct file *file;
};

struct proc {
    .....

    struct vma vma[VMASIZE];    // vma of file system
}
```

- sysfile里实现mmap，设置参数

```
uint64 sys_mmap(void)
{
    uint64 addr;
    int length, prot, flags, fd, offset;
    struct file *file;
    struct proc *p = myproc();

    // 读参数
    if (argaddr(0, &addr) || argint(1, &length) || argint(2, &prot) ||
        argint(3, &flags) || argfd(4, &fd, &file) || argint(5, &offset))
    {
        return -1;
    }

    // 读写权限判定
    if (!file->writable && (prot & PROT_WRITE) && flags == MAP_SHARED)
        return -1;

    // 检测有没有超出长度
```

```

length = PGROUNDUP(length);
if (p->sz > MAXVA - length)
    return -1;

// 找到第一个空vma
for (int i = 0; i < VMASIZE; i++)
{
    if (p->vma[i].used == 0)
    {
        p->vma[i].used = 1;
        p->vma[i].addr = p->sz;
        p->vma[i].length = length;
        p->vma[i].prot = prot;
        p->vma[i].flags = flags;
        p->vma[i].fd = fd;
        p->vma[i].file = file;
        p->vma[i].offset = offset;
        filedup(file);
        p->sz += length;
        return p->vma[i].addr;
    }
}
return -1;
}

```

- 这个实验要求的也是lazy allocated，因此和cow一样，在trap里处理映射导致的读写权限冲突的问题。检查到传入的va没有问题后就给它申请一个新的vma

```

else if (r_scause() == 13 || r_scause() == 15)
{ // 13是读缺页中断，15是写缺页中断
    // 检查是否超过页大小，是否超过地址，是否越页访问
    uint64 va = r_stval();
    if (va >= p->sz || va > MAXVA || PGROUNDUP(va) == PGROUNDDOWN(p->trapframe->sp))
        p->killed = 1;
    else
    {
        struct vma *vma = 0;
        for (int i = 0; i < VMASIZE; i++)
        {
            if (p->vma[i].used == 1 && va >= p->vma[i].addr && va < p->vma[i].addr + p->vma[i].length)
            {
                vma = &p->vma[i];
                break;
            }
        }
        if (vma)
        {
            va = PGROUNDDOWN(va);
            uint64 offset = va - vma->addr;
            uint64 mem = (uint64)kalloc();
            if (mem == 0)
            {

```

```

        p->killed = 1;
    }
    else
    {
        memset((void *)mem, 0, PGSIZE);
        ilock(vma->file->ip);
        readi(vma->file->ip, 0, mem, offset, PGSIZE);
        iunlock(vma->file->ip);
        int flag = PTE_U;
        if (vma->prot & PROT_READ)
            flag |= PTE_R;
        if (vma->prot & PROT_WRITE)
            flag |= PTE_W;
        if (vma->prot & PROT_EXEC)
            flag |= PTE_X;
        if (mmap_pages(p->pagetable, va, PGSIZE, mem, flag) != 0)
        {
            kfree((void *)mem);
            p->killed = 1;
        }
    }
}
}
}
}

```

- lazy allocate 中，当设置虚拟内存和物理内存映射时，可以先不复制虚拟内存；只有实际访问内存时才复制，因此有可能PTE\_V位可以0，只有当访问后触发中断复制才设置为1。因此需要去掉uvmunmap和uvmcopy的panic
- 取消掉指定的映射

```

uint64 sys_munmap(void)
{
    uint64 addr;
    int length;
    struct proc *p = myproc();
    struct vma *vma = 0;
    if (argaddr(0, &addr) || argint(1, &length))
        return -1;

    addr = PGROUNDDOWN(addr);
    length = PGROUNDUP(length);

    // 遍历进程的虚拟内存区域列表，查找包含给定地址的VMA
    for (int i = 0; i < VMASIZE; i++)
    {
        if (addr >= p->vma[i].addr || addr < p->vma[i].addr + p->vma[i].length)
        {
            vma = &p->vma[i];
            break;
        }
    }
    if (vma == 0)

```

```

    return 0;

    // 如果找到了起始地址
    if (vma->addr == addr)
    {
        vma->addr += length;
        vma->length -= length;

        // 如果vma是共享映射就更新数据
        if (vma->flags & MAP_SHARED)
            fwrite(vma->file, addr, length);
        uvmunmap(p->pagetable, addr, length / PGSIZE, 1);
    }

    // 如果VMA的长度为0，表示已经取消映射完毕
    if (vma->length == 0)
    {
        fclose(vma->file);
        vma->used = 0;
    }
    return 0;
}

```

- 最后，fork需要把映射复制过去，exit需要去掉映射

fork:

```

// 遍历，把vma复制一遍
for (int i = 0; i < VMASIZE; i++)
{
    if (p->vma[i].used)
    {
        memmove(&(np->vma[i]), &(p->vma[i]), sizeof(p->vma[i]));
        filedup(p->vma[i].file);
    }
}

```

exit:

```

for (int i = 0; i < VMASIZE; i++)
{
    if (p->vma[i].used)
    {
        if (p->vma[i].flags & MAP_SHARED)
            fwrite(p->vma[i].file, p->vma[i].addr, p->vma[i].length);
        fclose(p->vma[i].file);
        uvmunmap(p->pagetable, p->vma[i].addr, p->vma[i].length / PGSIZE, 1);
        p->vma[i].used = 0;
    }
}

```

## 实验结果

---

```
== Test running mmaptest == (5.6s)
== Test    mmaptest: mmap f ==
    mmaptest: mmap f: OK
== Test    mmaptest: mmap private ==
    mmaptest: mmap private: OK
== Test    mmaptest: mmap read-only ==
    mmaptest: mmap read-only: OK
== Test    mmaptest: mmap read/write ==
    mmaptest: mmap read/write: OK
== Test    mmaptest: mmap dirty ==
    mmaptest: mmap dirty: OK
== Test    mmaptest: not-mapped unmap ==
    mmaptest: not-mapped unmap: OK
== Test    mmaptest: two files ==
    mmaptest: two files: OK
== Test    mmaptest: fork_test ==
    mmaptest: fork_test: OK
== Test usertests == usertests: OK (135.5s)
== Test time ==
time: FAIL
    Cannot read time.txt
Score: 139/140
```

## 实验小结

---

又是被lazy allocated折磨的一个实验

难点主要在于对vma理解 (va-vma-pa)，现在还是理解为类似于用户态-文件态的一个过渡态页表（

每个进程在其虚拟地址空间中都有多个 VMA，每个 VMA 表示一段连续的虚拟地址范围，它可以包含程序的代码段、数据段、堆、栈等。操作系统使用 VMA 来管理虚拟内存的分配、映射和保护。）

然后是各种copy的报错

写完啦！